

Neural Networks Part 3

Kevin Sim



Slides courtesy of Prof. Emma Hart

RECAP

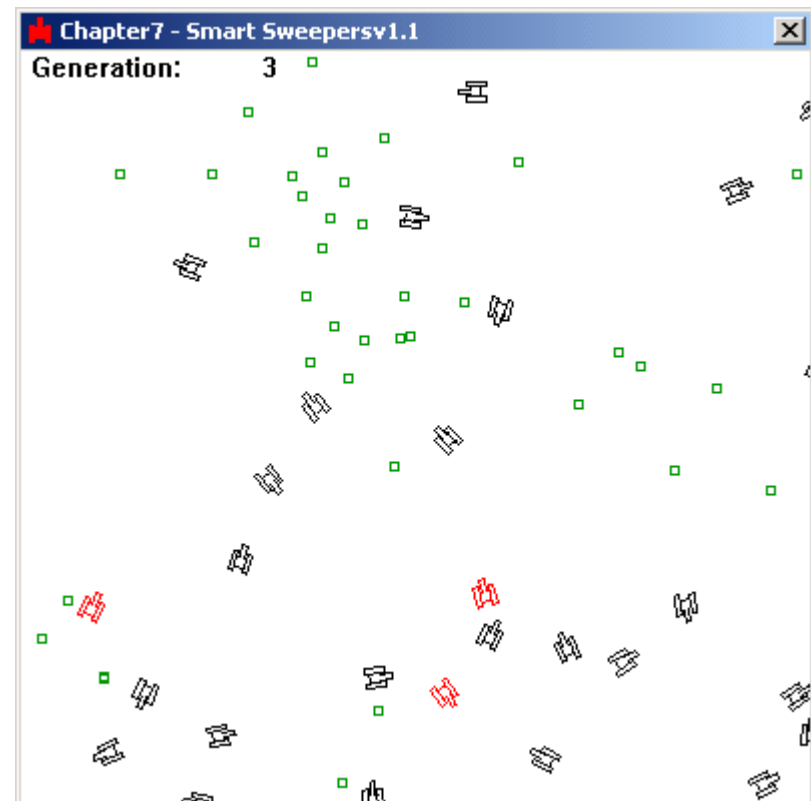
- Last week we looked at supervised learning using a multi-layer perceptron for classification
- Backpropagation
 - Finds a set of weights that gives good performance
 - Requires training data: input and output pairs
 - Iteratively reduces error at the output neurons
- For some applications:
 - Hard to generate the required training data (pairs)
 - We don't know what the output should be

Overview

- Unsupervised learning using NN
- Training a neural network with an evolutionary algorithm
 - i.e. use EA to find a set of weights
- Applications of this:
 - Robot control with a neural network
- NEAT: evolving topology & weights of a network with an EA

An Example: Minesweeper

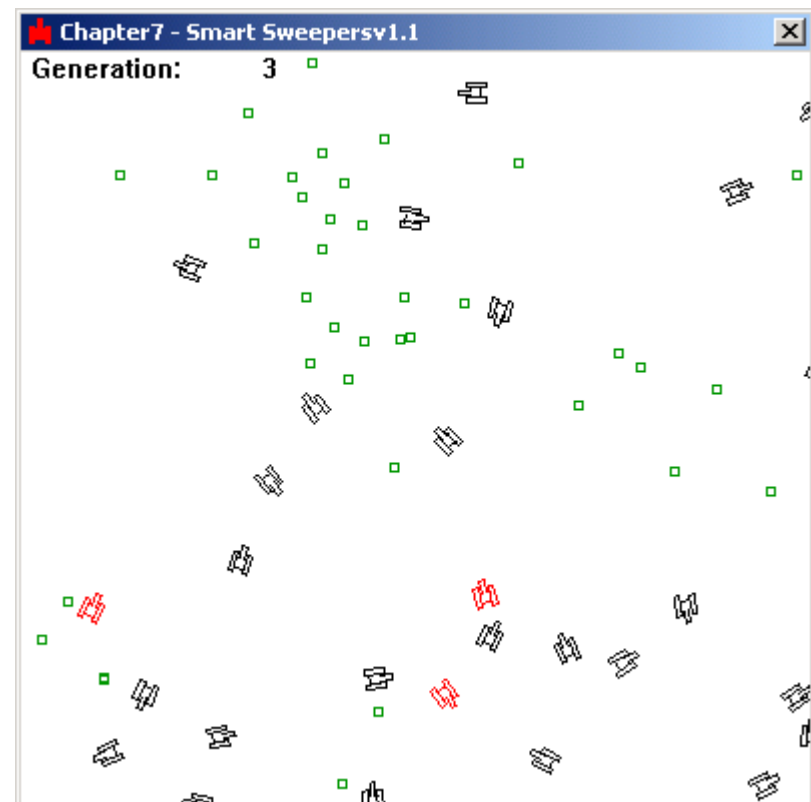
- Objective is to collect mines
- Neural network controls the movement of a “robot” (minesweeper)
- From last week we know a NN requires inputs and outputs
- Inputs are sensory information obtained from the environment
- Output is direction of movement
- NN Controller is trained to direct the minesweeper towards mines




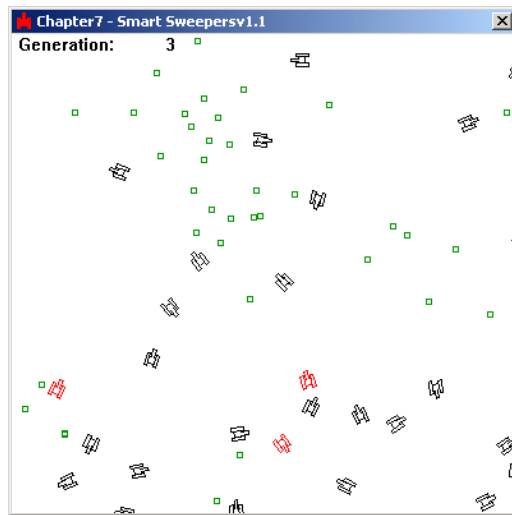
An Example: Minesweeper

- Objective is to collect mines
- EA used to evolve weights for the NN controller
- EA Requires a representation (chromosome genotype) and fitness function (more later)
- Example: Natural Motion

<https://www.youtube.com/watch?v=ySRvKzZsDqw>




Minesweeper: Inputs



- In robotics, inputs obtained by sensing the environment
- In real-robots, could be infra-red, acoustic, video etc.
- In simulation, we can 'sense' distance to an obstacle

Minesweeper: Inputs

- What information does it need ?
 - Minesweeper's position (x,y)
 - Position of closest mine (x,y)
 - Vector representing the heading $(x,y)^*$
- **6 inputs in total**




Assume vectors relative to origin so defined by 2 values

Minesweeper: Inputs

- Could use 6 inputs:


But

- Fewer inputs: fewer weights
- Fewer weights: faster training
- Faster training: faster network




Minesweeper: Inputs

- Actual positions don't matter
- The relative direction between the mine and the current direction **does**
- Input can actually be represented by **2 vectors**
- This only requires 4 inputs:
 - (x,y) for each vector




Minesweeper: Inputs

- NN require normalised data
- Good idea to standardize inputs:
 - Look-at vector is normalised to be of length 1
 - Closest-mine vector might be very large
 - We can normalize this too
- Both inputs then have similar emphasis
 - Actual distances and directions aren't important just relative difference between look at and closest mine




Minesweeper: Outputs

- The rotation and velocity are adjusted by activating one or both of the left and right tracks
- The NN needs two outputs:
 - Left track
 - Right track
- Rotate by moving one track




Minesweeper: Outputs

- The rotation and velocity are adjusted by activating one or both of the left and right tracks
- The NN needs two outputs:
 - Left track
 - Right track
- Move forward by activating both tracks




Digression:

- When we talked about simple perceptrons we used a step function that output 0/1
- More useful to have a neuron that can output any value
 - Probability
 - Distance
 - Angle





Threshold Functions

- Rather than a step activation function, it would be better to have one that varied:
 - Smoothly
 - Continuously
- Why ?
 - Can output a whole range of values
 - There is no abrupt change from one value to another
 - Imagine a car engine control system that allowed only 0 or maximum power



Minesweeper: Outputs


- The rotation and velocity are adjusted by adjusting the **relative** speed of the left and right tracks
- The NN needs two outputs:
 - Left track speed
 - Right track speed
 - Smooth transition from 0 (off) to 1 (maximum)
 - Allows the robot to turn while moving position



Threshold Functions

- There are lots of mathematical functions that might work
- This one is called a **sigmoid function**
 - Varies between 0 and 1
 - Shape can be 'flattened' if required
 - It is define as:

$$\frac{1}{1 + e^{-\frac{a}{p}}}$$




The Sigmoid Function

- a is the weighted sum


$$\sum w_i x_i$$

- e is a mathematical constant 2.7183
- p is a user controlled parameter usually set to 1
- (by changing p we can 'squash' the curve)

$$\frac{1}{1 + e^{-\frac{a}{p}}}$$



The Minesweeper Neural Net



We need to assign each weight so that the network outputs the correct values for any set of possible inputs

Calculating the weights

- In this case, there are 72 weights to find
- For backpropagation, we need training data with input-output pairs:
 - Hard to obtain for this type of application:
- Another approach to training is to use an **Evolutionary Algorithm** to evolve the weights
 - A chromosome (length= number of weights) represents the weights in neural network

w_1 w_2 w_3

.....

w_{71} w_{72}

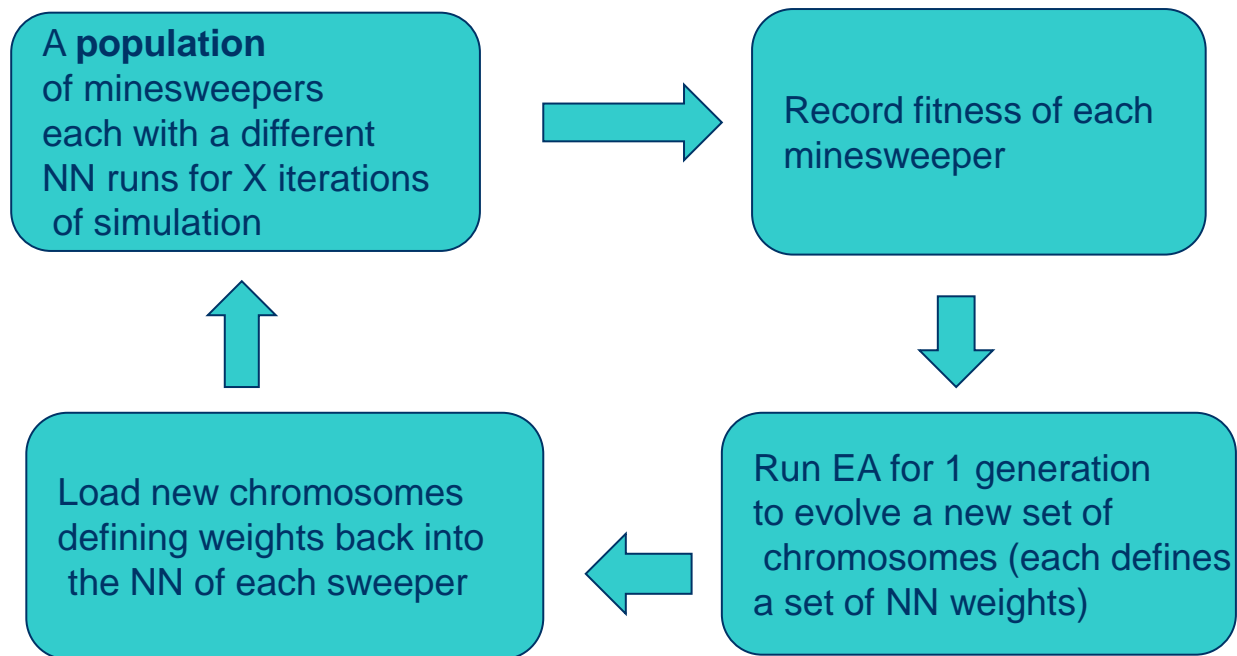
In-Game Training with an EA

- An EA needs a **representation**
 - Use floating point values, one for each weight we need
- and a **fitness function**:
 - Allow the game to run for some number of frames
 - Monitor how many mines each sweeper found

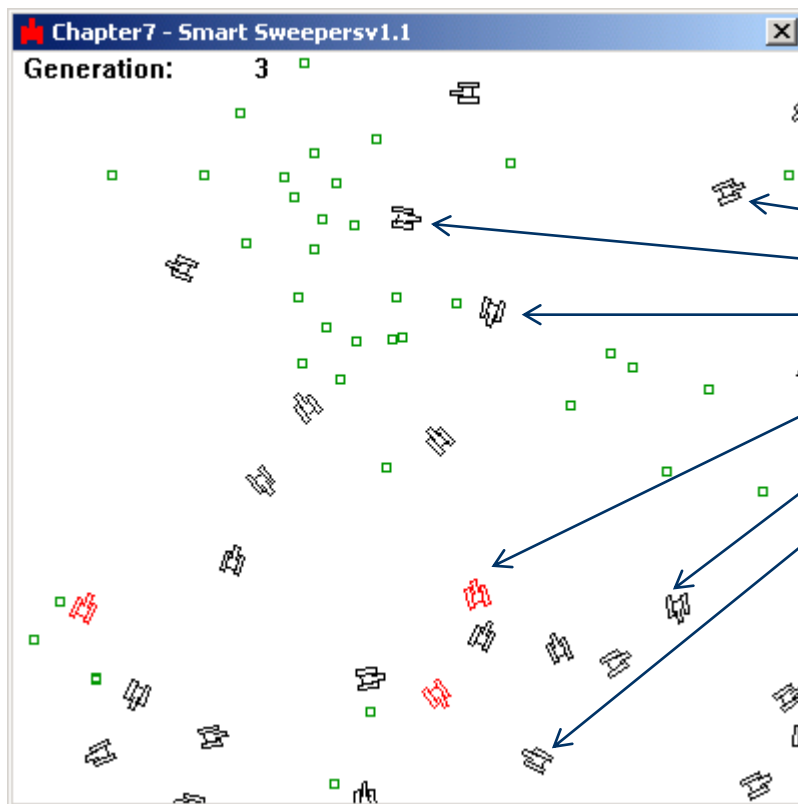
-0.1	0.3	0.6	0.4	-0.8
------	-----	-----	-----	-------	------

**Fitness = number of
mines found in fixed
time frame**

The Flow of Control



Training with an EA



Chromosome 1 (NN1)
Chromosome 2 (NN2)
Chromosome 3 (NN3)
0.1 0.3 0.5...0.3

EA
population

Chromosome represents
weights of a neural network

Flow of Control

Initialise a population of minesweepers each with its own neural network and its own chromosome representing the weights

Weights initially set to random values

While number of generations < maximum generations

- Run game (for a set amount of time)
- Record how many mines each sweeper detected
- Assign this value as the fitness of its chromosome
- Run the EA to evolve a new population of weights
 - Apply selection, crossover, mutation
- Insert new weights into each minesweepers NN

Repeat

Evolutionary Algorithm

- Generational EA (whole new generation produced each iteration)
- Differs from the steady state EA where typically 1 or 2 children are generated each generation

BEST Chromosome 1
Chromosome 2
Chromosome 3
.
.
.
.
Chromosome N

Generation 

Copy best chromosome to new population (**elitism**)

While newPopSize < popSize
Generate child

- Selection
- Crossover
- Mutation

Insert Child (**no replacement**)

Repeat


BEST Chromosome 1
Child Chromosome 2
Child Chromosome 3
.
.
.
.
Child Chromosome N

Evolutionary Algorithm

- In each generation:
 - Create new empty population
 - Copy best chromosome from current to new population (elitism)
 - Repeat until new population full:
 - Select two parents with **roulette wheel selection** from current population
 - Apply **crossover** to produce new child(ren)
 - Apply **mutation** to new child(ren)
 - Add child(ren) to new population
 - New population becomes current population


More details on the EA

- Selection:
 - Selection pressure is as a result of the selection stage only
 - No replacement stage in a generational EA
 - Best chromosome (could be more than 1. i.e. best 2) copied to new population each generation (elitism)
- Common to use Roulette Selection
 - Selection probability proportional to fitness



More details on the EA

- Crossover:
 - Any kind of crossover will produce legitimate chromosomes (1pt etc.)



- Two point or uniform crossover will work just as well

More details on the EA

- Mutation
 - Adds or subtracts a small value from each gene

```
for (int i=0;i<chromoSize;i++){  
    if (randFloat < mutationRate)  
        chromo[i] += randomClamped()*maxPerturbation;  
}
```

Probability of
mutation

Random
number
between
-1 and 1

Fix max size of
change

Summary of EA+NN

Neural Network

- Inputs: 4
- Outputs: 2
- Hidden Layers: 1
- Hidden Neurons: 10


Evolutionary Algorithm

- Population Size: 30
- Selection Type: roulette
- Crossover type: 1 point
- Mutation Rate: 0.1
- Elitism: on
- Max Perturbation 0.3

Run for approx. 2000 generations to fully
train networks

Some Improvements:

- We can reduce number of inputs even further
 - The important information is the **angle** between the two vectors
 - If we calculate the angle, we can just have one input to the network
 - **How ?**



Inputs

- We can calculate the angle using the **dot product**


$$A.B = |A| |B| \cos\theta$$

$$A.B = a_x b_x + a_y b_y$$


$$a_1 b_1 + a_2 b_2 = |A| |B| \cos\theta$$

$$a_1 b_1 + a_2 b_2 = \cos\theta \quad (\text{if magnitude normalised to 1})$$

- We also need to know if the **relative position** of the mine to the heading (left or right): use a sign +/-, e.g.
 - 30 = 30 degrees to left
 - +15 = 15 degrees to right




More improvements



- Weights are listed in the chromosome *per neuron*
- 1pt (or 2pt) crossover makes a random cut
- This can break up the weights for a single neuron
- Better to choose crossover points that only occur at the boundaries of neurons
 - So this limits the number of cut points


Obstacle Avoidance with NNs

- Obstacle avoidance common requirement in game AI
 - while still exploring the environment
- To do this successfully, agents need to:
 - perceive environment
 - take action to avoid collisions
- Typically implemented by adding sensors to the robot/car/agent etc....




Obstacle Avoidance with NNs

- We can add sensors to each agent
- Adjust the number and range as required
- Each sensor returns:
 - -1 if there is no intersection between it and an obstacle
 - a value between 0 and 1 indicating the distance to the intersection otherwise
 - (closer to 0, closer to object)







Evolving an NN Controller

- Inputs to the NN are the readings from the 5 sensors
- Outputs again are the speeds of the left/right tracks
- Try one hidden layer at first
- Use an EA to evolve the correct weights



The Evolutionary Algorithm

- The EA needs a fitness function:
- What is a 'good' controller ?
 - Could record # collisions and penalise everytime it collides
 - Higher fitness = better minesweeper 
 - Could lead to negative scores 
- Better:
 - Record number of frames passed without colliding
 - More frames, higher fitness 
 - Doesn't lead to negative scores 

The Fitness Function

- Does the fitness function look reasonable ?
 - Fitness: #frames passed without collision

```
if (!collided) fitness++;
```

The Fitness Function

- Does the fitness function look reasonable ?
 - Fitness: #frames passed without collision
- Robot will just rotate on the spot!!

```
if (!collided) fitness++;
```

The Fitness Function

- Does the fitness function look reasonable ?
 - Fitness: #frames passed without collision
- Robot will just rotate on the spot!!
- How can we fix this ?
 - Add extra reward for frames where there is zero (or little) rotation

```
if (!collided) fitness++;  
if (rotation < rotationTolerance) fitness++;
```

Fitness Function

- Goal is to maximise fitness
- For each frame that passes:


```
if (!Collided){  
    fitness +=1  
}  
  
if (abs(Rotation) < RotationTolerance){  
    fitness += 1  
}
```

Amount robot rotated

Threshold defining max allowed rotation


Further improvements

- We want the robot to explore the environment
- Add memory:
 - Record how many times each square is visited
 - Evolve networks that favour unvisited cells
- The sensors can 'feel' how many times a square has been visited



Further improvements

- Readings from sensors converted to scaled value between 0 and 1 representing times visited
- Sliding scale important to give robot a sense of 'direction'



Fitness Function:

- We could combine the previous function with an extra factor:

$\text{Fitness} = \text{NoRotation} + \text{NoCollision} + \text{Cells visited}$

- But....can just use number of cells visited:
 - Automatically will avoid obstacles
 - Automatically will stop spinning
 - (both slow them down and reduce fitness)

Uses in “real games”

- Calculate aiming for the AI characters in Quake ?
 - To prevent them being too accurate (...and therefore unrealistic)
- Network inputs:
 - Visibility (dark, foggy, bright)
 - Amount of target visible
 - Distance to target
 - Current weapon
- Output
 - Radius of distance from centre of target (bigger radius = less accurate)



Some additional thoughts

- In order to satisfy the fitness function many robot behaviours can emerge
 - Robot driving example
- You can also train a network to classify data using an EA
 - Fitness function = total error

Evolving Neural Network Topology


- As well as evolving weights, it makes sense to evolve topology as well
- A very well known technique is called NEAT:
 - Neuro-Evolution of Augmenting Topologies
 - Starts by using a population of networks with minimal technology that grow in complexity
- Generates neurons, connections, layers and weights

<http://www.cs.ucf.edu/~kstanley/neat.html>

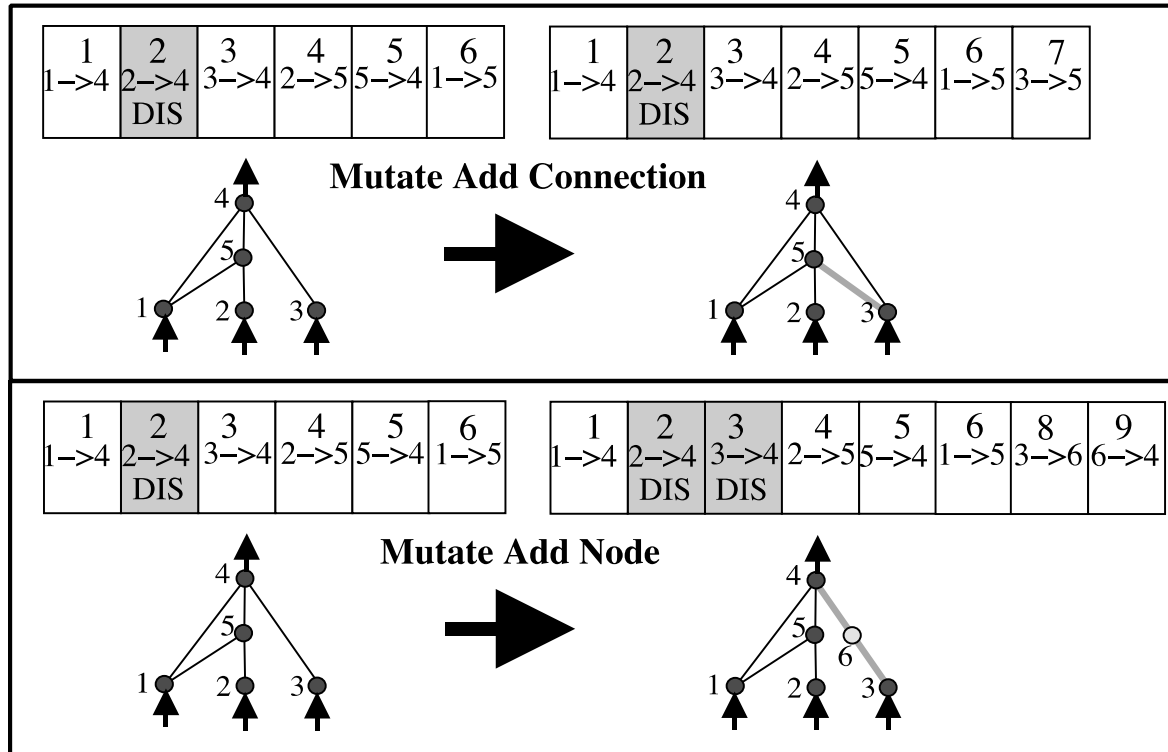
NEAT genomes

Genome (Genotype)							
Node Genes	Node						
	Node 1	Node 2	Node 3	Node 4	Node 5		
Connect. Genes	Sensor	Sensor	Sensor	Output	Hidden		
	In 1 Out 4 Weight 0.7 Enabled Innov 1	In 2 Out 4 Weight -0.5 DISABLED Innov 2	In 3 Out 4 Weight 0.5 Enabled Innov 3	In 2 Out 5 Weight 0.2 Enabled Innov 4	In 5 Out 4 Weight 0.4 Enabled Innov 5	In 1 Out 5 Weight 0.6 Enabled Innov 6	In 4 Out 5 Weight 0.6 Enabled Innov 11





Network (Phenotype)



NEAT Mutation operators



Example: NERO



[Video](#)
[robot driving](#)

Comparison of Training Methods

Backpropagation

- Data needs to have input/output pairs
- Need plenty of data
- Careful data cleaning
- Training and testing required
- Lots of software packages available

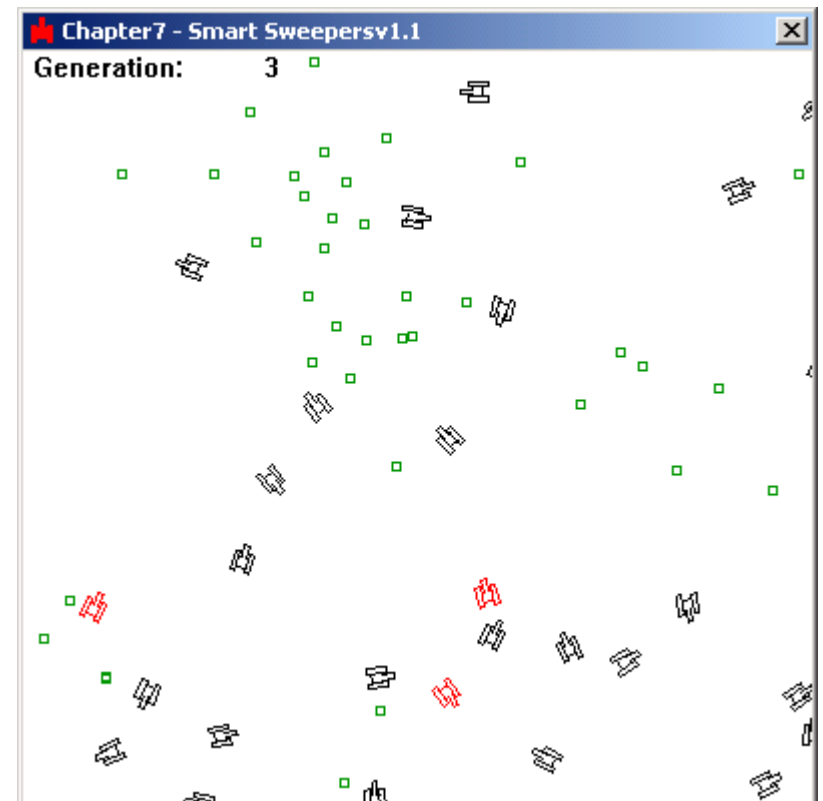
Evolutionary Algorithm

- Good when you don't have training data in the form of input/outputs pairs
- Careful design of fitness functions required
- Lots of parameters: need to deal with EA parameters as well as NN design
- Can be slow....

Practical

Evolving Minesweeper controller

- Source Code and Executable supplied with initialisation file
- Adapt topology and parameters by changing values in initialisation file
 - Conduct experiments to evaluate effect
- Implement own evolutionary operators by modifying C++ Code
 - Selection
 - Mutation
 - Crossover
- Other possibilities
 - Try different activation functions



Summary

- We have looked at a number of uses of Neural Networks:
 - Classification
 - Prediction
 - Control
- Training a network can be performed by backpropagation or using an evolutionary algorithm
 - Stochastic process in both cases
- A lot of effort goes into cleaning and preparing data
- Testing with unseen data very important